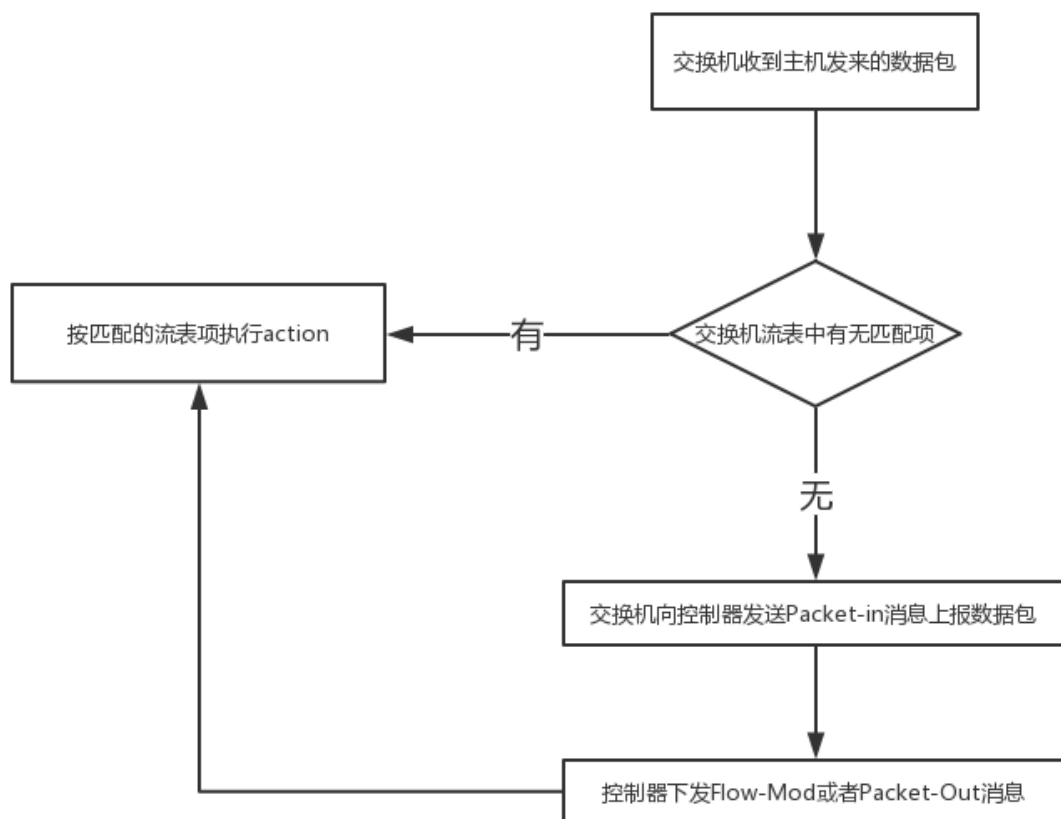
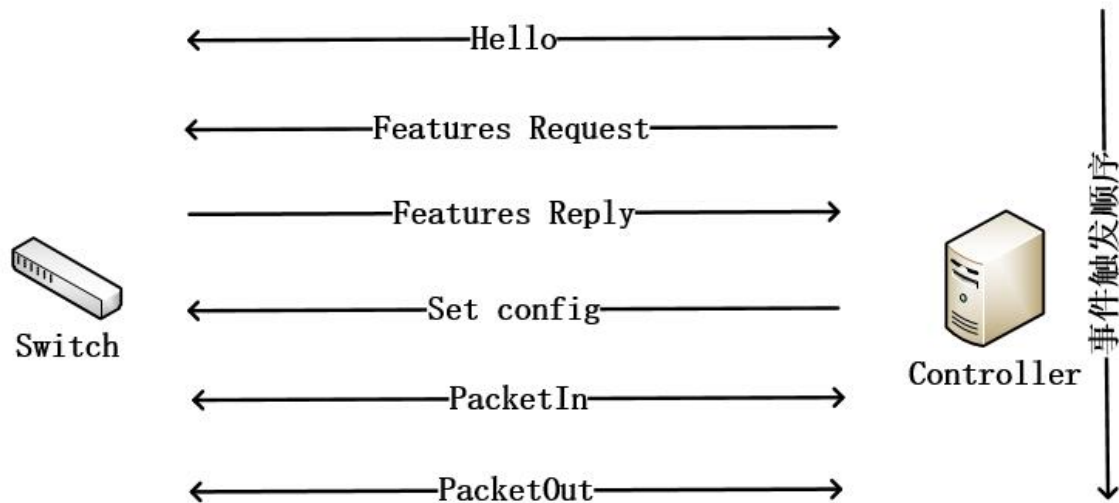


Lab: RYU Programming

1. The Step of Openflow switch interacting with controller



Assignments

Topo

Requirements

Set up the following network first:

```
      s3
    /   \
h1 - s1   s2 - h2
    \   /
      s4
```

```
1  #!/usr/bin/python
2  from mininet.topo import Topo
3  from mininet.net import Mininet
4  from mininet.log import setLogLevel
5  from mininet.cli import CLI
6  from mininet.node import OVSSwitch, Controller, RemoteController
7  from time import sleep
8
9
10 class SingleSwitchTopo(Topo):
11     "Single switch connected to n hosts."
12     def build(self):
13         s1 = self.addSwitch('s1')
14         s2 = self.addSwitch('s2')
15         s3 = self.addSwitch('s3')
16         s4 = self.addSwitch('s4')
17
18         h1 = self.addHost('h1')
19         h2 = self.addHost('h2')
20         h1 = self.addHost('h1')
21         h2 = self.addHost('h2')
22         self.addLink(s1, s4, 1, 1)
23         self.addLink(s1, s3, 2, 1)
24         self.addLink(s4, s2, 2, 2)
25         self.addLink(s3, s2, 2, 1)
26         self.addLink(s1, h1, 3, 1)
27         self.addLink(s2, h2, 3, 1)
28
29 if __name__ == '__main__':
30     setLogLevel('info')
31     topo = SingleSwitchTopo()
32     c1 = RemoteController('c1', ip='127.0.0.1')
33     net = Mininet(topo=topo, controller=c1)
34     net.start()
35     CLI(net)
36     net.stop()
```

auto switch per 5secs

Requirements

Write a RYU controller that switches paths (h1-s1-s3-s2-h2 or h1-s1-s4-s2-h2) between h1 and h2 every 5 seconds.

Ideas

1. set a `hard_timeout` flow.
2. after 5 seconds the timer is out, it will trigger the event
`ofp_event.EventOFPFlowRemoved`
3. set a `_flow_removed_handler` to add new flow with another 5secs `hard_timeout` to the data path.

Code

```
1  #http://docs.openvswitch.org/en/latest/faq/openflow/
2  #https://mail.openvswitch.org/pipermail/ovs-discuss/2016-August/042394.html
3  #https://stackoverflow.com/questions/36949861/group-table-issue-openflow-
   mininet
4
5  from ryu.base import app_manager
6  from ryu.controller import ofp_event
7  from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
8  from ryu.controller.handler import set_ev_cls
9  from ryu.ofproto import ofproto_v1_3
10 from ryu.lib.packet import packet
11 from ryu.lib.packet import ethernet
12 from ryu.lib.packet import ether_types
13 s1op = 1
14 s2op = 1
15
16 class SimpleSwitch13(app_manager.RyuApp):
17     OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
18
19     def __init__(self, *args, **kwargs):
20         super(SimpleSwitch13, self).__init__(*args, **kwargs)
21         self.mac_to_port = {}
22         def add_flow_tmp(self, datapath, priority, match, actions,
buffer_id=None):
23             ofproto = datapath.ofproto
24             parser = datapath.ofproto_parser
25
26             inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
27                                                 actions)]
28             if buffer_id:
29                 mod = parser.OFPFlowMod(datapath=datapath,
buffer_id=buffer_id, hard_timeout=5,
30
31                 flags=ofproto.OFPFF_SEND_FLOW_REM, priority=priority, match=match,
28                                     instructions=inst)
32             else:
33                 mod = parser.OFPFlowMod(datapath=datapath,
34                 priority=priority, hard_timeout=5,
35
36                 flags=ofproto.OFPFF_SEND_FLOW_REM, match=match, instructions=inst)
```

```

35     datapath.send_msg(mod)
36
37 @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
38 def switch_features_handler(self, ev):
39     global s1op
40     global s2op
41     datapath = ev.msg.datapath
42     ofproto = datapath.ofproto
43     parser = datapath.ofproto_parser
44
45     # install table-miss flow entry
46     '''
47     match = parser.OFPMatch()
48     actions = [parser.OFPActionOutput(ofproto.OFPP_CONTROLLER,
49                                     ofproto.OFPCML_NO_BUFFER)]
50     self.add_flow(datapath, 0, match, actions)
51     '''
52     # switch s1
53     if datapath.id == 1:
54
55         #add the return flow for h1 in s1.
56         # h1 is connected to port 3.
57         actions = [parser.OFPActionOutput(s1op)]
58         match = parser.OFPMatch(in_port=3)
59         self.add_flow_tmp(datapath, 10, match, actions)
60         op = 2
61         actions = [parser.OFPActionOutput(3)]
62         match = parser.OFPMatch(in_port=1)
63         self.add_flow(datapath, 10, match, actions)
64
65         actions = [parser.OFPActionOutput(3)]
66         match = parser.OFPMatch(in_port=2)
67         self.add_flow(datapath, 10, match, actions)
68
69
70     # switch s2
71     if datapath.id == 2:
72         actions = [parser.OFPActionOutput(s2op)]
73         match = parser.OFPMatch(in_port=3)
74         self.add_flow_tmp(datapath, 10, match, actions)
75         op = 2
76         #add the return flow for h2 in s4.
77         # h2 is connected to port 3.
78         actions = [parser.OFPActionOutput(3)]
79         match = parser.OFPMatch(in_port=1)
80         self.add_flow(datapath, 10, match, actions)
81
82         actions = [parser.OFPActionOutput(3)]
83         match = parser.OFPMatch(in_port=2)
84         self.add_flow(datapath, 10, match, actions)
85
86
87     # switch s4
88     if datapath.id == 4:
89
90         actions = [parser.OFPActionOutput(2)]
91         match = parser.OFPMatch(in_port=1)
92         self.add_flow(datapath, 10, match, actions)

```

```

93         actions = [parser.OFPActionOutput(1)]
94         match = parser.OFPMatch(in_port=2)
95         self.add_flow(datapath, 10, match, actions)
96
97
98
99     # switch s3
100    if datapath.id == 3:
101        # h1 is connected to port 3.
102        actions = [parser.OFPActionOutput(2)]
103        match = parser.OFPMatch(in_port=1)
104        self.add_flow(datapath, 10, match, actions)
105
106        actions = [parser.OFPActionOutput(1)]
107        match = parser.OFPMatch(in_port=2)
108        self.add_flow(datapath, 10, match, actions)
109
110    def add_flow(self, datapath, priority, match, actions, buffer_id=None):
111        ofproto = datapath.ofproto
112        parser = datapath.ofproto_parser
113
114        inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
115                                             actions)]
116
117        if buffer_id:
118            mod = parser.OFPFlowMod(datapath=datapath, buffer_id=buffer_id,
119                                    priority=priority, match=match,
120                                    instructions=inst)
121        else:
122            mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
123                                    match=match, instructions=inst)
124        datapath.send_msg(mod)
125
126    @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
127    def _packet_in_handler(self, ev):
128        # If you hit this you might want to increase
129        # the "miss_send_length" of your switch
130        if ev.msg.msg_len < ev.msg.total_len:
131            self.logger.debug("packet truncated: only %s of %s bytes",
132                              ev.msg.msg_len, ev.msg.total_len)
133
134        msg = ev.msg
135        datapath = msg.datapath
136        ofproto = datapath.ofproto
137        parser = datapath.ofproto_parser
138        in_port = msg.match['in_port']
139
140        pkt = packet.Packet(msg.data)
141        eth = pkt.get_protocols(ethernet.ethernet)[0]
142
143        if eth.ethertype == ether_types.ETH_TYPE_LLDP:
144            # ignore lldp packet
145            return
146
147        dst = eth.dst
148        src = eth.src
149
150        dpid = datapath.id
151        if dst[:5] == "33:3":
152            self.logger.info("drop ipv6 multicast packet %s", dst)

```

```

151         return
152
153     self.mac_to_port.setdefault(dpid, {})
154
155     self.logger.info("packet in %s %s %s %s", dpid, src, dst, in_port)
156
157     if(dpid==1 or dpid == 4):
158         if(op==1):
159             out_port=2
160             op = 2
161         else:
162             out_port = 1
163             op = 1
164         actions = [parser.OFPActionOutput(out_port)]
165
166         # install a flow to avoid packet_in next time
167         if out_port != ofproto.OFPP_FLOOD:
168             match = parser.OFPMatch(in_port=in_port, eth_dst=dst,
eth_src=src)
169             # verify if we have a valid buffer_id, if yes avoid to send both
170             # flow_mod & packet_out
171             if msg.buffer_id != ofproto.OFP_NO_BUFFER:
172                 self.add_flow_tmp(datapath, 1, match, actions,
msg.buffer_id)
173             return
174         else:
175             self.add_flow_tmp(datapath, 1, match, actions)
176         data = None
177         if msg.buffer_id == ofproto.OFP_NO_BUFFER:
178             data = msg.data
179
180         out = parser.OFPPacketOut(datapath=datapath,
buffer_id=msg.buffer_id,
181                                 in_port=in_port, actions=actions,
data=data)
182         datapath.send_msg(out)
183
184     @set_ev_cls(ofp_event.EventOFPPFlowRemoved, MAIN_DISPATCHER)
185     def _flow_removed_handler(self, ev):
186         global s1op
187         global s2op
188         msg = ev.msg
189         dp = msg.datapath
190         ofp = dp.ofproto
191         parser = dp.ofproto_parser
192         if msg.reason == ofp.OFPRR_IDLE_TIMEOUT:
193             reason = 'IDLE TIMEOUT'
194         elif msg.reason == ofp.OFPRR_HARD_TIMEOUT:
195             reason = 'HARD TIMEOUT'
196         elif msg.reason == ofp.OFPRR_DELETE:
197             reason = 'DELETE'
198         elif msg.reason == ofp.OFPRR_GROUP_DELETE:
199             reason = 'GROUP DELETE'
200         else:
201             reason = 'unknown'
202
203         self.logger.debug('OFPPFlowRemoved received: ' 'datapath id=%s'
204                         'cookie=%d priority=%d reason=%s table_id=%d '

```

```

205         'duration_sec=%d duration_nsec=%d '
206         'idle_timeout=%d hard_timeout=%d '
207         'packet_count=%d byte_count=%d
match.fields=%s',
208         dp.id,msg.cookie, msg.priority, reason,
msg.table_id,
209         msg.duration_sec, msg.duration_nsec,
210         msg.idle_timeout, msg.hard_timeout,
211         msg.packet_count, msg.byte_count, msg.match)
212     if(dp.id==1):
213         if(slop==1):
214             slop = 2
215         else:
216             slop = 1
217     print("slop is "+str(slop))
218     actions = [parser.OFPActionOutput(slop)]
219     match = parser.OFPMatch(in_port=3)
220     self.add_flow_tmp(dp, 10, match, actions)

```

Test and Result

1. start the controller

```
ryu-manager --verbose per5s.py
```

2. start net topo and connect to the controller

```
sudo python topo.py --controller remote
```

In mininet using `h1 ping h2` test the connections. Using `sudo ovs-ofctl -O OpenFlow13 dump-flows s1` check the flow table.

Mininet Result

```

mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.991 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.088 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.089 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.087 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.665 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.089 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=0.095 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=0.089 ms
64 bytes from 10.0.0.2: icmp_seq=9 ttl=64 time=0.092 ms
64 bytes from 10.0.0.2: icmp_seq=10 ttl=64 time=0.090 ms
64 bytes from 10.0.0.2: icmp_seq=11 ttl=64 time=0.090 ms
64 bytes from 10.0.0.2: icmp_seq=12 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=13 ttl=64 time=0.090 ms
64 bytes from 10.0.0.2: icmp_seq=14 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=15 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=16 ttl=64 time=0.093 ms
64 bytes from 10.0.0.2: icmp_seq=17 ttl=64 time=0.090 ms
64 bytes from 10.0.0.2: icmp_seq=18 ttl=64 time=0.089 ms
^C
--- 10.0.0.2 ping statistics ---
18 packets transmitted, 18 received, 0% packet loss, time 17393ms
rtt min/avg/max/mdev = 0.087/0.172/0.991/0.238 ms

```

RYU Controller Result:

```

EVENT ofp_event->SimpleSwitch13 EventOFFFlowRemoved
OFFFlowRemoved received: datapath id=1cookie=0 priority=10 reason=HARD TIMEOUT table_id=0 duration_sec=5 duration_nsec=16000000 idle_timeout=0 hard_timeo
t=5 packet_count=5 byte_count=490 match.fields=OFFMatch(oxm_fields={'in_port': 3})
stop is 1
EVENT ofp_event->SimpleSwitch13 EventOFFFlowRemoved
OFFFlowRemoved received: datapath id=2cookie=0 priority=10 reason=HARD TIMEOUT table_id=0 duration_sec=5 duration_nsec=18000000 idle_timeout=0 hard_timeo
t=5 packet_count=5 byte_count=490 match.fields=OFFMatch(oxm_fields={'in_port': 3})
stop is 1
EVENT ofp_event->SimpleSwitch13 EventOFFFlowRemoved
OFFFlowRemoved received: datapath id=1cookie=0 priority=10 reason=HARD TIMEOUT table_id=0 duration_sec=5 duration_nsec=17000000 idle_timeout=0 hard_timeo
t=5 packet_count=5 byte_count=490 match.fields=OFFMatch(oxm_fields={'in_port': 3})
stop is 2
EVENT ofp_event->SimpleSwitch13 EventOFFFlowRemoved
OFFFlowRemoved received: datapath id=2cookie=0 priority=10 reason=HARD TIMEOUT table_id=0 duration_sec=5 duration_nsec=19000000 idle_timeout=0 hard_timeo
t=5 packet_count=6 byte_count=532 match.fields=OFFMatch(oxm_fields={'in_port': 3})
stop is 2

```

It show the event triggered the handler and made echo.

Flow table:

```

austinguish@austinguish-GL502VSK ~/桌/C/C/L/r/r/app (master)> sudo ovs-ofctl -O OpenFlow13 dump-flows s2
cookie=0x0, duration=23.995s, table=0, n_packets=83, n_bytes=11974, priority=10,in_port="s2-eth1" actions=output:"s2-eth3"
cookie=0x0, duration=23.995s, table=0, n_packets=122, n_bytes=15604, priority=10,in_port="s2-eth2" actions=output:"s2-eth3"
cookie=0x0, duration=3.919s, table=0, n_packets=3, n_bytes=294, hard_timeout=5, send_flow_rem priority=10,in_port="s2-eth3" actions=output:"s2-eth1"
austinguish@austinguish-GL502VSK ~/桌/C/C/L/r/r/app (master)> sudo ovs-ofctl -O OpenFlow13 dump-flows s2
cookie=0x0, duration=25.225s, table=0, n_packets=83, n_bytes=11974, priority=10,in_port="s2-eth1" actions=output:"s2-eth3"
cookie=0x0, duration=25.225s, table=0, n_packets=124, n_bytes=15800, priority=10,in_port="s2-eth2" actions=output:"s2-eth3"
cookie=0x0, duration=0.129s, table=0, n_packets=0, n_bytes=0, hard_timeout=5, send_flow_rem priority=10,in_port="s2-eth3" actions=output:"s2-eth2"

```

Load Balancer

Requirements

Write a RYU controller that uses both paths to forward packets from h1 to h2.

Ideas:

1. set the basic flow in `switch_features_handler`
2. set a group table to handle the forwarding.
3. It will have a weight to select action in buckets.
4. using iperf -P set multi connection

Code

```

1  from ryu.base import app_manager
2  from ryu.controller import ofp_event
3  from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
4  from ryu.controller.handler import set_ev_cls
5  from ryu.ofproto import ofproto_v1_3
6  from ryu.lib.packet import packet
7  from ryu.lib.packet import ethernet
8  from ryu.lib.packet import ether_types
9
10 class SimpleSwitch13(app_manager.RyuApp):
11     OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
12
13     def __init__(self, *args, **kwargs):
14         super(SimpleSwitch13, self).__init__(*args, **kwargs)
15         self.mac_to_port = {}
16
17     @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
18     def switch_features_handler(self, ev):
19         datapath = ev.msg.datapath
20         ofproto = datapath.ofproto
21         parser = datapath.ofproto_parser
22         if datapath.id == 1:
23             # add group tables
24             self.send_group_mod(datapath)
25             actions = [parser.OFPActionGroup(group_id=50)]
26             match = parser.OFPMatch(in_port=3)

```



```

27         self.add_flow(datapath, 10, match, actions)
28
29         #add the return flow for h1 in s1.
30         # h1 is connected to port 3.
31         actions = [parser.OFPActionOutput(3)]
32         match = parser.OFPMatch(in_port=1)
33         self.add_flow(datapath, 10, match, actions)
34
35         actions = [parser.OFPActionOutput(3)]
36         match = parser.OFPMatch(in_port=2)
37         self.add_flow(datapath, 10, match, actions)
38
39
40     # switch s2
41     if datapath.id == 2:
42     # add group tables
43         self.send_group_mod(datapath)
44         actions = [parser.OFPActionGroup(group_id=50)]
45         match = parser.OFPMatch(in_port=3)
46         self.add_flow(datapath, 10, match, actions)
47
48
49         #add the return flow for h2 in s2.
50         # h2 is connected to port 3.
51         actions = [parser.OFPActionOutput(3)]
52         match = parser.OFPMatch(in_port=1)
53         self.add_flow(datapath, 10, match, actions)
54
55         actions = [parser.OFPActionOutput(3)]
56         match = parser.OFPMatch(in_port=2)
57         self.add_flow(datapath, 10, match, actions)
58
59
60     # switch s4
61     if datapath.id == 4 :
62         actions = [parser.OFPActionOutput(2)]
63         match = parser.OFPMatch(in_port=1)
64         self.add_flow(datapath, 10, match, actions)
65
66         actions = [parser.OFPActionOutput(1)]
67         match = parser.OFPMatch(in_port=2)
68         self.add_flow(datapath, 10, match, actions)
69
70
71     # switch s3
72     if datapath.id == 3:
73         actions = [parser.OFPActionOutput(2)]
74         match = parser.OFPMatch(in_port=1)
75         self.add_flow(datapath, 10, match, actions)
76
77         actions = [parser.OFPActionOutput(1)]
78         match = parser.OFPMatch(in_port=2)
79         self.add_flow(datapath, 10, match, actions)
80
81     def add_flow(self, datapath, priority, match, actions, buffer_id=None):
82         ofproto = datapath.ofproto
83         parser = datapath.ofproto_parser
84

```

```

85         inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
86                                             actions)]
87     if buffer_id:
88         mod = parser.OFPFlowMod(datapath=datapath, buffer_id=buffer_id,
89                                 priority=priority, match=match,
90                                 instructions=inst)
91     else:
92         mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
93                                 match=match, instructions=inst)
94     datapath.send_msg(mod)
95
96     def send_group_mod(self, datapath):
97         ofproto = datapath.ofproto
98         parser = datapath.ofproto_parser
99         LB_WEIGHT1 = 50 #percentage
100        LB_WEIGHT2 = 50 #percentage
101
102        watch_port = ofproto_v1_3.OFPP_ANY
103        watch_group = ofproto_v1_3.OFPQ_ALL
104
105        actions1 = [parser.OFPActionOutput(1)]
106        actions2 = [parser.OFPActionOutput(2)]
107        buckets = [parser.OFPBucket(LB_WEIGHT1, watch_port, watch_group,
108                                    actions=actions1),
109                    parser.OFPBucket(LB_WEIGHT2, watch_port, watch_group,
110                                    actions=actions2)]
111        req = parser.OFPGroupMod(datapath, ofproto.OFPGC_ADD,
112                                ofproto.OFPGT_SELECT, 50, buckets)
113        datapath.send_msg(req)

```

Test and Result

```
mininet> h2 iperf -s &
```

```
mininet> h1 iperf -c h2 -P 4 -t 30
```

using `sudo ovs-ofctl -O OpenFlow13 dump-group-stats s1` watch the load

```

austinguish@austinguish-GL502VSK ~/C/C/L/r/r/app (master)> sudo ovs-ofctl -O OpenFlow13 dump-group-stats s1
OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=50,duration=732.090s,ref_count=1,packet_count=7465592,byte_count=346749069788,bucket0:packet_count=3707482,byte_count=17154847354
  0,bucket1:packet_count=3758110,byte_count=175200596240

```

bucket0: Total 3707482 packets

bucket1: Total 3758110 packets

```

austinguish@austinguish-GL502VSK ~/C/C/L/r/r/app (master)> sudo ovs-ofctl -O OpenFlow13 dump-group-stats s2
OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=50,duration=944.588s,ref_count=1,packet_count=5239935,byte_count=345863226,bucket0:packet_count=2652978,byte_count=175109476,buck
  et1:packet_count=2586957,byte_count=170753750

```

bucket0: Total 2652978 packets

bucket1: Total 2585957 packets

using `sudo ovs-ofctl -O OpenFlow13 dump-flows s3 s4` watch the load

```

austinguish@austinguish-GL502VSK ~/C/C/L/r/r/app (master)> sudo ovs-ofctl -O OpenFlow13 dump-flows s3
cookie=0x0, duration=834.669s, table=0, n_packets=3758142, n_bytes=175200601364, priority=10,in_port="s3-eth1" actions=output:"s3-eth2"
cookie=0x0, duration=834.669s, table=0, n_packets=2653011, n_bytes=175114686, priority=10,in_port="s3-eth2" actions=output:"s3-eth1"
austinguish@austinguish-GL502VSK ~/C/C/L/r/r/app (master)> sudo ovs-ofctl -O OpenFlow13 dump-flows s4
cookie=0x0, duration=842.645s, table=0, n_packets=3707514, n_bytes=171548478672, priority=10,in_port="s4-eth1" actions=output:"s4-eth2"
cookie=0x0, duration=842.645s, table=0, n_packets=2586990, n_bytes=170758960, priority=10,in_port="s4-eth2" actions=output:"s4-eth1"

```

The results show the load on 2 paths is balanced

Fast failover

Requirements

Write a RYU controller that uses the first path (h1-s1-s3-s2-h2) for routing packets from h1 to h2 and uses the second path for backup. Specifically, when the first path experiences a link failure, the network should automatically switch to the second path without causing packet drop. (hint: consider using OFPGT_FF (FF is short for "fast failover") to construct a group table)

Idea

1. To `switch 1` and `switch2` all packets from port 3 are forwarded to `group table 50`
2. Set `OFPGT_FF` group mode to monitor the port 1 and port 2.
3. Auto switch the port.

Code

```
1  from ryu.base import app_manager
2  from ryu.controller import ofp_event
3  from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
4  from ryu.controller.handler import set_ev_cls
5  from ryu.ofproto import ofproto_v1_3
6  from ryu.lib.packet import packet
7  from ryu.lib.packet import ethernet
8  from ryu.lib.packet import ether_types
9
10 class SimpleSwitch13(app_manager.RyuApp):
11     OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
12
13     def __init__(self, *args, **kwargs):
14         super(SimpleSwitch13, self).__init__(*args, **kwargs)
15         self.mac_to_port = {}
16
17     @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
18     def switch_features_handler(self, ev):
19         datapath = ev.msg.datapath
20         ofproto = datapath.ofproto
21         parser = datapath.ofproto_parser
22         if datapath.id == 1:
23             # add group tables
24             self.send_group_mod(datapath)
25             actions = [parser.OFPActionGroup(group_id=50)]
26             match = parser.OFPMatch(in_port=3)
27             self.add_flow(datapath, 10, match, actions)
28
29             #add the return flow for h1 in s1.
30             # h1 is connected to port 3.
31             actions = [parser.OFPActionOutput(3)]
32             match = parser.OFPMatch(in_port=1)
33             self.add_flow(datapath, 10, match, actions)
34
35             actions = [parser.OFPActionOutput(3)]
36             match = parser.OFPMatch(in_port=2)
37             self.add_flow(datapath, 10, match, actions)
38
39
40     # switch s2
```

```

41     if datapath.id == 2:
42         # add group tables
43         self.send_group_mod(datapath)
44         actions = [parser.OFPActionGroup(group_id=50)]
45         match = parser.OFPMatch(in_port=3)
46         self.add_flow(datapath, 10, match, actions)
47
48
49         #add the return flow for h2 in s2.
50         # h2 is connected to port 3.
51         actions = [parser.OFPActionOutput(3)]
52         match = parser.OFPMatch(in_port=1)
53         self.add_flow(datapath, 10, match, actions)
54
55         actions = [parser.OFPActionOutput(3)]
56         match = parser.OFPMatch(in_port=2)
57         self.add_flow(datapath, 10, match, actions)
58
59
60     # switch s4
61     if datapath.id == 4 :
62         actions = [parser.OFPActionOutput(2)]
63         match = parser.OFPMatch(in_port=1)
64         self.add_flow(datapath, 10, match, actions)
65
66         actions = [parser.OFPActionOutput(1)]
67         match = parser.OFPMatch(in_port=2)
68         self.add_flow(datapath, 10, match, actions)
69
70
71     # switch s3
72     if datapath.id == 3:
73         actions = [parser.OFPActionOutput(2)]
74         match = parser.OFPMatch(in_port=1)
75         self.add_flow(datapath, 10, match, actions)
76
77         actions = [parser.OFPActionOutput(1)]
78         match = parser.OFPMatch(in_port= 2)
79         self.add_flow(datapath, 10, match, actions)
80
81     def add_flow(self, datapath, priority, match, actions, buffer_id=None):
82         ofproto = datapath.ofproto
83         parser = datapath.ofproto_parser
84
85         inst = [parser.OFPInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
86                                             actions)]
87         if buffer_id:
88             mod = parser.OFPFlowMod(datapath=datapath, buffer_id=buffer_id,
89                                     priority=priority, match=match,
90                                     instructions=inst)
91         else:
92             mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
93                                     match=match, instructions=inst)
94         datapath.send_msg(mod)
95
96     def send_group_mod(self, datapath):
97         ofproto = datapath.ofproto
98         parser = datapath.ofproto_parser

```

```

99     actions1 = [parser.OFPActionOutput(1)]
100    actions2 = [parser.OFPActionOutput(2)]
101    buckets = [parser.OFPBucket(watch_port=1,actions=actions1),
102               parser.OFPBucket(watch_port=2,actions=actions2)]
103    req = parser.OFPGroupMod(datapath, ofproto.OFPGC_ADD,
104                             ofproto.OFPGT_FF, 50, buckets)
105    datapath.send_msg(req)

```

Test and Result

use `link s1 s3 down` `link s2 s3 down` to close the connection

```

64 bytes from 10.0.0.2: icmp_seq=27 ttl=64 time=0.087 ms
64 bytes from 10.0.0.2: icmp_seq=28 ttl=64 time=0.094 ms
64 bytes from 10.0.0.2: icmp_seq=29 ttl=64 time=0.089 ms
64 bytes from 10.0.0.2: icmp_seq=30 ttl=64 time=0.092 ms
64 bytes from 10.0.0.2: icmp_seq=31 ttl=64 time=0.055 ms
64 bytes from 10.0.0.2: icmp_seq=32 ttl=64 time=0.106 ms
64 bytes from 10.0.0.2: icmp_seq=33 ttl=64 time=0.085 ms
64 bytes from 10.0.0.2: icmp_seq=34 ttl=64 time=0.057 ms
64 bytes from 10.0.0.2: icmp_seq=35 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=36 ttl=64 time=0.054 ms
64 bytes from 10.0.0.2: icmp_seq=37 ttl=64 time=0.062 ms
64 bytes from 10.0.0.2: icmp_seq=38 ttl=64 time=0.050 ms
64 bytes from 10.0.0.2: icmp_seq=39 ttl=64 time=0.063 ms
64 bytes from 10.0.0.2: icmp_seq=40 ttl=64 time=0.103 ms
64 bytes from 10.0.0.2: icmp_seq=41 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=45 ttl=64 time=0.441 ms
64 bytes from 10.0.0.2: icmp_seq=46 ttl=64 time=0.098 ms
64 bytes from 10.0.0.2: icmp_seq=47 ttl=64 time=0.105 ms
64 bytes from 10.0.0.2: icmp_seq=48 ttl=64 time=0.094 ms
64 bytes from 10.0.0.2: icmp_seq=49 ttl=64 time=0.093 ms
64 bytes from 10.0.0.2: icmp_seq=50 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=51 ttl=64 time=0.089 ms
64 bytes from 10.0.0.2: icmp_seq=52 ttl=64 time=0.089 ms
64 bytes from 10.0.0.2: icmp_seq=53 ttl=64 time=0.045 ms
64 bytes from 10.0.0.2: icmp_seq=54 ttl=64 time=0.085 ms
64 bytes from 10.0.0.2: icmp_seq=55 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=56 ttl=64 time=0.090 ms
64 bytes from 10.0.0.2: icmp_seq=57 ttl=64 time=0.090 ms
64 bytes from 10.0.0.2: icmp_seq=58 ttl=64 time=0.086 ms

```

There are some packet loss between the icmp_seq42-44. Because when the connection closed between `s1` and `s3` the `s2` didn't know the path is loss, it still forwards packet to `s3`. The solution is to set a backup link between the `s3` and `s4`, set `fast failover table` in `s3` and `s4`.

Solution

new topo

```

1           Switch3
2           /   |   \
3  h1 ---Switch1 |   Switch2-----h2
4           \   |   /
5           Switch4
6
7

```

```

1 from ryu.base import app_manager
2 from ryu.controller import ofp_event
3 from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
4 from ryu.controller.handler import set_ev_cls

```

```

5 from ryu.ofproto import ofproto_v1_3
6 from ryu.lib.packet import packet
7 from ryu.lib.packet import ethernet
8 from ryu.lib.packet import ether_types
9
10 class SimpleSwitch13(app_manager.RyuApp):
11     OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
12
13     def __init__(self, *args, **kwargs):
14         super(SimpleSwitch13, self).__init__(*args, **kwargs)
15         self.mac_to_port = {}
16
17     @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
18     def switch_features_handler(self, ev):
19         datapath = ev.msg.datapath
20         ofproto = datapath.ofproto
21         parser = datapath.ofproto_parser
22         if datapath.id == 1:
23             # add group tables
24             self.send_group_mod(datapath)
25             actions = [parser.OFPActionGroup(group_id=50)]
26             match = parser.OFPMatch(in_port=3)
27             self.add_flow(datapath, 10, match, actions)
28
29             #add the return flow for h1 in s1.
30             # h1 is connected to port 3.
31             actions = [parser.OFPActionOutput(3)]
32             match = parser.OFPMatch(in_port=1)
33             self.add_flow(datapath, 10, match, actions)
34
35             actions = [parser.OFPActionOutput(3)]
36             match = parser.OFPMatch(in_port=2)
37             self.add_flow(datapath, 10, match, actions)
38
39
40         # switch s2
41         if datapath.id == 2:
42             # add group tables
43             self.send_group_mod(datapath)
44             actions = [parser.OFPActionGroup(group_id=50)]
45             match = parser.OFPMatch(in_port=3)
46             self.add_flow(datapath, 10, match, actions)
47
48
49             #add the return flow for h2 in s2.
50             # h2 is connected to port 3.
51             actions = [parser.OFPActionOutput(3)]
52             match = parser.OFPMatch(in_port=1)
53             self.add_flow(datapath, 10, match, actions)
54
55             actions = [parser.OFPActionOutput(3)]
56             match = parser.OFPMatch(in_port=2)
57             self.add_flow(datapath, 10, match, actions)
58
59
60         # switch s4
61         if datapath.id == 4 :
62             actions = [parser.OFPActionOutput(2)]

```


I never set the weight in the topo, so the default path `h1` to `h2` is `h1-s1-s4-s2-h2` .

The back route is `h2-s2-s3-s1-h1`

When I haven't cut the link between `s1` `s3` the result.

```

OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=51,duration=238.773s,ref_count=1,packet_count=94,byte_count=10798,bucket0:packet_count=94,byte_count=10798,bucket1:packet_count=0,byte_count=0
austinguish@austinguish-GL502VSK ~> sudo ovs-ofctl -O OpenFlow13 dump-group-stats s3
OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=51,duration=240.722s,ref_count=1,packet_count=96,byte_count=10994,bucket0:packet_count=96,byte_count=10994,bucket1:packet_count=0,byte_count=0
austinguish@austinguish-GL502VSK ~> sudo ovs-ofctl -O OpenFlow13 dump-group-stats s3
OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=51,duration=255.665s,ref_count=1,packet_count=111,byte_count=12436,bucket0:packet_count=111,byte_count=12436,bucket1:packet_count=0,byte_count=0

```

The bucket 1 is never used

After I cut the link between `s1` `s3` the result:

```

austinguish@austinguish-GL502VSK ~> sudo ovs-ofctl -O OpenFlow13 dump-group-stats s3
OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=51,duration=375.946s,ref_count=1,packet_count=238,byte_count=24471,bucket0:packet_count=224,byte_count=23099,bucket1:packet_count=14,byte_count=1372
austinguish@austinguish-GL502VSK ~> sudo ovs-ofctl -O OpenFlow13 dump-group-stats s3
OFPST_GROUP reply (OF1.3) (xid=0x6):
  group_id=51,duration=376.924s,ref_count=1,packet_count=239,byte_count=24569,bucket0:packet_count=224,byte_count=23099,bucket1:packet_count=15,byte_count=1470

```

And there is no packet loss:

```

64 bytes from 10.0.0.2: icmp_seq=214 ttl=64 time=0.091 ms
64 bytes from 10.0.0.2: icmp_seq=215 ttl=64 time=0.051 ms
64 bytes from 10.0.0.2: icmp_seq=216 ttl=64 time=0.108 ms
64 bytes from 10.0.0.2: icmp_seq=217 ttl=64 time=0.099 ms
64 bytes from 10.0.0.2: icmp_seq=218 ttl=64 time=0.040 ms
64 bytes from 10.0.0.2: icmp_seq=219 ttl=64 time=0.100 ms
64 bytes from 10.0.0.2: icmp_seq=220 ttl=64 time=0.049 ms
64 bytes from 10.0.0.2: icmp_seq=221 ttl=64 time=0.086 ms
64 bytes from 10.0.0.2: icmp_seq=222 ttl=64 time=0.096 ms
64 bytes from 10.0.0.2: icmp_seq=223 ttl=64 time=0.092 ms
^C
--- 10.0.0.2 ping statistics ---
223 packets transmitted, 223 received, 0% packet loss, time 227243ms
rtt min/avg/max/mdev = 0.036/0.093/1.670/0.110 ms

```